

Aside: related topics and extensions

- The **Recurrent Neural Network Transducer** (RNN-T) extends CTC by allowing dependencies between b_t and $a_{1,\dots,u-1}$
 - $table[u, t] = P(b_t = a_u | a_{1,\dots,u-1}, x)(table[u-1, t-1] + table[u, t-1])$
 - Compute p_u with auto-regressive RNN, then combine with p_t
- Partial derivatives w.r.t. the loss \mathcal{L} can be efficiently calculated with the **forward_backward** algorithm
 - **monotonic_backward** builds suffixes instead of prefixes
 - Multiplies $\frac{\delta P_{p_t}(b_t=a_u)}{\delta p_t}$ with prefix and suffix likelihoods
- Setting $b = x$ and a as a state sequence, **forward_backward** can be used to train a **GMM-HMM**
 - $P_u(b_t)(a_{u-1,u}table[u-1, t-1] + a_{u,u}table[u, t-1])$
- More details can be found in the appendices

Everything past here is an **ASIDE**

(not on your exam)

Backpropagating the loss 1

- To learn $\mathcal{L} = -\log P_\theta(a)$ end-to-end, we need to **backprop**
- The parameters for frame t are p_t , and

$$\frac{\partial \mathcal{L}}{\partial p_t} = -\frac{1}{P_\theta(a)} \frac{\partial P_\theta(a)}{\partial p_t} = -\frac{1}{P_\theta(a)} \sum_{b \in \mathcal{B}^{-1}(a; T)} \frac{\partial P_\theta(b|x)}{\partial p_t}$$

- For a single alignment b we have

$$\begin{aligned} \frac{\partial P_\theta(b|x)}{\partial p_t} &= \frac{\partial \prod_{t'=1}^T P_{p_{t'}}(b_{t'})}{\partial p_t} \\ &= \left(\prod_{t'=1}^{t-1} P_{p_{t'}}(b_{t'}) \right) \left(\prod_{t'=t+1}^T P_{p_{t'}}(b_{t'}) \right) \frac{\partial P_{p_t}(b_t)}{\partial p_t} \\ &= P_\theta(b_{1,\dots,t-1}|x) P_\theta(b_{t+1,\dots,T}|x) \frac{\partial P_{p_t}(b_t)}{\partial p_t} \end{aligned}$$

Backpropagating the loss 2

- Considering **all** paths $\mathcal{B}(b) = a$, we have

$$\begin{aligned}
 \frac{\partial \mathcal{L}}{\partial p_t} &= -\frac{1}{P_\theta(a)} \sum_{b \in \mathcal{B}^{-1}(a;T)} \frac{\partial P_\theta(b|x)}{\partial p_t} \\
 &= -\frac{1}{P_\theta(a)} \sum_{b \in \mathcal{B}^{-1}(a;T)} P_\theta(b_{1,\dots,t-1}|x) P_\theta(b_{t+1,\dots,T}|x) \frac{\partial P_{p_t}(b_t)}{\partial p_t} \\
 &= -\frac{1}{P_\theta(a)} \sum_{u=1}^U \frac{\partial P_{p_t}(b_t = a_u)}{\partial p_t} \left(\sum_{b_{1,\dots,t-1} \in \mathcal{B}^{-1}(a_{1,\dots,u-1};t-1)} P_\theta(b_{1,\dots,t-1}|x) \right) \quad \leftarrow \textcircled{1} \\
 &\quad \left(\sum_{b_{t+1,\dots,T} \in \mathcal{B}^{-1}(a_{u+1,\dots,U};T-t-1)} P_\theta(b_{t+1,\dots,T}|x) \right) \quad \leftarrow \textcircled{2}
 \end{aligned}$$

- $\textcircled{1}$ is $table[u-1, t-1]$ in `monotonic_forward`
- $\textcircled{2}$ is $table[u+1, t+1]$ in `monotonic_backward`

The monotonic_backward algorithm

Function monotonic_backward

Inputs $a = a_1, a_2, \dots, a_U$ and $b = b_1, b_2, \dots, b_T$

1: **Define** $table[1 \dots U + 1, 1 \dots T + 1]$

2: initialize($table[1 \dots U + 1, T + 1], table[U + 1, 1 \dots T + 1]$)

3: **For each** u in $U \dots 1$:

4: **For each** t in $T \dots 1$:


5: $table[u, t] =$

$step(a_u, b_t, table[u + 1, t + 1], table[u + 1, t], table[u, t + 1])$

6: **Return** finalize($table[1 \dots U + 1, 1 \dots T + 1]$)

Same as monotonic_forward, but with partial solutions over **suffixes**

The forward_backward algorithm



- Prefixes are computed with the monotonic_forward algorithm (➡)
- Suffixes are computed with the monotonic_backward algorithm (⬅)
- $\frac{\partial}{\partial p_t}$ are summed over columns (a_u)


HMM-based ASR

- In GMM-based phone classification, $P(x_t|q_t)$ depends on which phone $q_t = s$ we assume x_t was drawn from
- We can model temporal dependencies across frames by specifying $P(q)$
 - $q^* = \operatorname{argmax}_q P(q)P(x|q)$ is a frame-level transcript
- Using a **Hidden Markov Model** (HMM) over GMM-based observation likelihoods, we have a GMM-HMM ASR system

$$P(q, x) = \underbrace{P(q_o)}_{\text{Prior state probability}} \prod_{t=1}^T \underbrace{P(q_t|q_{t-1})}_{\text{Transition probability}} \underbrace{P(x_t|q_t)}_{\text{Observation probability (GMM)}}$$

Continuous HMMs (CHMM)


- A **continuous HMM** has observations that are distributed over continuous variables.
 - Observation probabilities, P_s , are also continuous.
 - E.g., here $b_s(x_t)$ tells us the probability of seeing the (multivariate) continuous observation x_t while in state s .



GMM-HMM as Continuous HMM

- Continuous HMMs are very similar to discrete HMMs.
 - $S = \{s_1, \dots, s_V\}$: set of states (e.g., phonemes)
 - $X = \mathbb{R}^d$: **continuous observation space**

$$\theta \left\{ \begin{array}{l} \bullet \Pi = \{\pi_1, \dots, \pi_V\} \\ \bullet A = \{a_{ij}\}, i, j \in S \\ \bullet B = P_v(\vec{x}), i \in S, \vec{x} \in X \end{array} \right. \begin{array}{l} \text{: initial state probabilities} \\ \text{: state transition probabilities} \\ \text{: **state output probabilities**} \\ \text{(**i.e., Gaussian mixtures**)}$$



yielding

- $Q = \{q_1, \dots, q_T\}, q_t \in S$: state sequence
- $\mathcal{O} = \{\sigma_1, \dots, \sigma_T\}, \sigma_t \in X$: observation sequence

Adapting monotonic_forward

Function monotonic_forward

Inputs $a = a_1, a_2, \dots, a_U$ and $b = b_1, b_2, \dots, b_T$

1: **Define** $table[0 \dots U, 0 \dots T]$

2: initialize($table[0 \dots U, 0], table[0, 1 \dots T]$)

3: **For each** u in $1 \dots U$:

4: **For each** t in $1 \dots T$:

5: $table[u, t] =$
 $step(a_u, b_t, table[u - 1, t - 1], table[u - 1, t], table[u, t - 1])$

6: **Return** finalize($table[0 \dots U, 0 \dots T]$)

a is the state sequence, $b = x$ the speech features

initialize: set $table[0, 0] = \pi_{q_0}$, $table[0, 1 \dots T] = table[1 \dots U, 0] = 0$

step: $P_u(b_t)(a_{u-1,u}table[u-1, t-1] + a_{u,u}table[u, t-1])$

finalize: $-\log table[U, T]$

(Though training usually involves EM, not backprop)